

# BIG DATA LAB HIVE ASSIGNMENT 1

---

**Name : Asha Venugopal Madguni**

**USN : 1NT18IS033**

**Sem and Sec : 6-'A' (A2 Batch)**

- **Create database called customer**

```
hive> create database Customer;
OK
Time taken: 5.329 seconds
hive> use Customer;
OK
Time taken: 0.13 seconds
hive> create table bank(tid int,bname string,cname string,amount float)row format delimited fields terminated by ",";
OK
Time taken: 3.487 seconds
hive> insert into bank values(101,"SBI","Asha",50000.00),(102,"Axis","Ananya",30000.00),(103,"BOB","Goutan",80000.00),(104,"Canara","Nandita",60000.00),(105,"KOC", "Pavitra",40000.00);
Query ID = hdoop_20210707123059_e0e8be0a-ed84-4746-8264-0dfa82d52639
Total jobs = 3
Launching Job 1 out of 3
Number of reduce tasks determined at compile time: 1
In order to change the average load for a reducer (in bytes):
  get hlog size, reduce size, see reducer number
```

- **Create table bank (tid, bname, cname, amount)**

```
hive> create table bank(tid int,bname string,cname string,amount float)row format delimited fields terminated by ",";
OK
Time taken: 3.487 seconds
hive> insert into bank values(101,"SBI","Asha",50000.00),(102,"Axis","Ananya",30000.00),(103,"BOB","Goutan",80000.00),(104,"Canara","Nandita",60000.00),(105,"KOC", "Pavitra",40000.00);
Query ID = hdoop_20210707123059_e0e8be0a-ed84-4746-8264-0dfa82d52639
Total jobs = 3
Launching Job 1 out of 3
Number of reduce tasks determined at compile time: 1
In order to change the average load for a reducer (in bytes):
  get hlog size, reduce size, see reducer number
```

- Insert values using insert statement

```

hadoop@asha-VirtualBox: ~/apache-hive-3.12-bin/conf
hive> select * from bank;
OK
101 SBI Asha 50000.0
102 Axis Ananya 30000.0
Time taken: 1.51 seconds, Fetched: 2 row(s)
hive> insert into bank values(103,"BOB","Goutam",80000.00),(104,"Canara","Nandita",60000.00),(105,"KDCB","Pavitra",40000.00);
Query ID = hdoop_20210707130917_Qec24a6b-eb8c-4144-8e05-43bbe242db04
Total jobs = 3
Launching Job 1 out of 3
Number of reduce tasks determined at compile time: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reducers=<number>
Starting Job = job_1625638408330_0003, Tracking URL = http://asha-VirtualBox:8088/proxy/application_1625638408330_0003/
Kill Command = /home/hadoop/hadoop-3.2.1/bin/mapred job -kill job_1625638408330_0003
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2021-07-07 13:10:14,561 Stage-1 map = 0%, reduce = 0%
2021-07-07 13:11:11,915 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 6.51 sec
2021-07-07 13:12:11,958 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 6.51 sec
2021-07-07 13:12:24,836 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 11.02 sec
MapReduce Total cumulative CPU time: 11 seconds 20 msec
Ended Job = job_1625638408330_0003
Stage-4 is selected by condition resolver.
Stage-3 is filtered out by condition resolver.
Stage-5 is filtered out by condition resolver.
Moving data to directory hdfs://127.0.0.1:9000/user/hive/warehouse/customer.db/bank/.hive-staging_hive_2021-07-07_13-09-17_577_7972249460389112203-1/-ext-10000
Loading data to table customer.bank
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 11.02 sec HDFS Read: 19974 HDFS Write: 466 SUCCESS
Total MapReduce CPU Time Spent: 11 seconds 20 msec
OK
Time taken: 201.75 seconds
hive> select * from bank;
OK
101 SBI Asha 50000.0
102 Axis Ananya 30000.0
103 BOB Goutam 80000.0
104 Canara Nandita 60000.0
105 KDCB Pavitra 40000.0
Time taken: 1.298 seconds, Fetched: 5 row(s)
hive> load data local inpath '/home/hadoop/customer.csv' into table bank;
Loading data to table customer.bank
OK
Time taken: 4.122 seconds
hive> select * from bank;
OK
101 SBI Asha 50000.0
102 Axis Ananya 30000.0
103 BOB Goutam 80000.0
104 Canara Nandita 60000.0
105 KDCB Pavitra 40000.0
106 Urban Venugopal 30000.0
107 SBI Bharati 55000.0
108 BOB Sindhu 65000.0
Time taken: 1.158 seconds, Fetched: 8 row(s)
hive> select count(*) from bank where amount < 40000 and name like '%SBI%';

```

- Import the values using csv

```

105 KDCB Pavitra 40000.0
Time taken: 1.298 seconds, Fetched: 5 row(s)
hive> load data local inpath '/home/hadoop/customer.csv' into table bank;
Loading data to table customer.bank
OK
Time taken: 4.122 seconds
hive> select * from bank;
OK
101 SBI Asha 50000.0
102 Axis Ananya 30000.0
103 BOB Goutam 80000.0
104 Canara Nandita 60000.0
105 KDCB Pavitra 40000.0
106 Urban Venugopal 30000.0
107 SBI Bharati 55000.0
108 BOB Sindhu 65000.0
Time taken: 1.158 seconds, Fetched: 8 row(s)
hive> select count(*) from bank where amount < 40000 and name like '%SBI%';

```

- Query using group by clause

```

hadoop@asha-VirtualBox: ~/apache-hive-3.1.2-bin/conf
104 Canara Nandita 60000.0
105 KDCB Pavitra 40000.0
Time taken: 1.298 seconds, Fetched: 5 row(s)
hive> load data local inpath '/home/hadoop/customer.csv' into table bank;
Loading data to table customer.bank
OK
Time taken: 4.122 seconds
hive> select * from bank;
OK
101 SBI Asha 50000.0
102 Axis Ananya 30000.0
103 BOB Goutam 80000.0
104 Canara Nandita 60000.0
105 KDCB Pavitra 40000.0
106 Urban Venugopal 30000.0
107 SBI Bharati 55000.0
108 BOB Sindhu 65000.0
Time taken: 1.158 seconds, Fetched: 8 row(s)
hive> select cname,amount from bank where amount>=50000.00 group by cname,amount;
Query ID = hadoop_20210707131535_2c4fe9ff-2ade-4c02-bcbe-0a60a1e6ea47
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1625638408330_0004, Tracking URL = http://asha-VirtualBox:8080/proxy/application_1625638408330_0004/
Kill Command = /home/hadoop/hadoop-3.2.1/bin/mapred job -kill job_1625638408330_0004
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2021-07-07 13:16:55,087 Stage-1 map = 0%, reduce = 0%
2021-07-07 13:17:34,494 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 6.07 sec
2021-07-07 13:18:19,514 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 10.58 sec
MapReduce Total cumulative CPU time: 10 seconds 910 msec
Ended Job = job_1625638408330_0004
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 10.91 sec HDFS Read: 13153 HDFS Write: 222 SUCCESS
Total MapReduce CPU Time Spent: 10 seconds 910 msec
OK
Asha 50000.0
Bharati 55000.0
Goutam 80000.0
Nandita 60000.0
Sindhu 65000.0
Time taken: 176.004 seconds, Fetched: 5 row(s)
hive> select tid,cname,amount from bank where amount>50000 order by amount DESC;
Query ID = hadoop_20210707132012_233cc0b8-0437-4cb0-bfa7-c421fc2005cd
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks determined at compile time: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1625638408330_0005, Tracking URL = http://asha-VirtualBox:8080/proxy/application_1625638408330_0005/
Kill Command = /home/hadoop/hadoop-3.2.1/bin/mapred job -kill job_1625638408330_0005
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2021-07-07 13:21:14,254 Stage-1 map = 0%, reduce = 0%
2021-07-07 13:22:03,577 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 5.88 sec
2021-07-07 13:23:04,296 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 5.88 sec
2021-07-07 13:23:08,538 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 9.68 sec
MapReduce Total cumulative CPU time: 9 seconds 680 msec
Ended Job = job_1625638408330_0005
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 9.68 sec HDFS Read: 12597 HDFS Write: 213 SUCCESS
Total MapReduce CPU Time Spent: 9 seconds 680 msec
OK
103 Goutam 80000.0
108 Sindhu 65000.0
104 Nandita 60000.0
107 Bharati 55000.0
Time taken: 185.854 seconds, Fetched: 4 row(s)

```

- Query using order by clause

```

Time taken: 176.004 seconds, Fetched: 5 row(s)
hive> select tid,cname,amount from bank where amount>50000 order by amount DESC;
Query ID = hadoop_20210707132012_233cc0b8-0437-4cb0-bfa7-c421fc2005cd
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks determined at compile time: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1625638408330_0005, Tracking URL = http://asha-VirtualBox:8080/proxy/application_1625638408330_0005/
Kill Command = /home/hadoop/hadoop-3.2.1/bin/mapred job -kill job_1625638408330_0005
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2021-07-07 13:21:14,254 Stage-1 map = 0%, reduce = 0%
2021-07-07 13:22:03,577 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 5.88 sec
2021-07-07 13:23:04,296 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 5.88 sec
2021-07-07 13:23:08,538 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 9.68 sec
MapReduce Total cumulative CPU time: 9 seconds 680 msec
Ended Job = job_1625638408330_0005
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 9.68 sec HDFS Read: 12597 HDFS Write: 213 SUCCESS
Total MapReduce CPU Time Spent: 9 seconds 680 msec
OK
103 Goutam 80000.0
108 Sindhu 65000.0
104 Nandita 60000.0
107 Bharati 55000.0
Time taken: 185.854 seconds, Fetched: 4 row(s)

```

- Query using having clause

```

Time taken: 185.854 seconds, Fetched: 4 row(s)
hive> select tid,cname,sum(amount) from bank group by tid,cname having sum(amount)>60000;
Query ID = hdoop_20210707132606_5ec869a4-a0cd-4317-bf6b-667e54bd131b
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1625638408330_0006, Tracking URL = http://asha-VirtualBox:8088/proxy/application_1625638408330_0006/
Kill Command = /home/hadoop/hadoop-3.2.1/bin/mapred job -kill job_1625638408330_0006
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2021-07-07 13:26:53,668 Stage-1 map = 0%, reduce = 0%
2021-07-07 13:27:30,953 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 4.59 sec
2021-07-07 13:28:30,035 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 10.23 sec
MapReduce Total cumulative CPU time: 10 seconds 230 msec
Ended Job = job_1625638408330_0006
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 10.23 sec HDFS Read: 14852 HDFS Write: 149 SUCCESS
Total MapReduce CPU Time Spent: 10 seconds 230 msec
OK
103 Goutam 80000.0
108 Sindhu 65000.0
Time taken: 148.852 seconds, Fetched: 2 row(s)
hive>

```

- Query on aggregate functions

```

hive> select max(amount) as MaxAmt,min(amount) as MinAmt,avg(amount) as AvgAmt;
FAILED: SemanticException [Error 10004]: Line 1:11 Invalid table alias or column reference 'amount': (possible column names are: )
hive> select max(amount) as MaxAmt,min(amount) as MinAmt,avg(amount) as AvgAmt from bank;
Query ID = hdoop_20210707133437_b7f61457-97ff-4825-aa2e-6a6e1f567806
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks determined at compile time: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1625638408330_0007, Tracking URL = http://asha-VirtualBox:8088/proxy/application_1625638408330_0007/
Kill Command = /home/hadoop/hadoop-3.2.1/bin/mapred job -kill job_1625638408330_0007
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2021-07-07 13:35:04,497 Stage-1 map = 0%, reduce = 0%
2021-07-07 13:36:04,657 Stage-1 map = 0%, reduce = 0%
2021-07-07 13:36:20,790 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 4.79 sec
2021-07-07 13:36:50,858 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 10.00 sec
MapReduce Total cumulative CPU time: 10 seconds 80 msec
Ended Job = job_1625638408330_0007
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 10.00 sec HDFS Read: 17237 HDFS Write: 123 SUCCESS
Total MapReduce CPU Time Spent: 10 seconds 80 msec
OK
80000.0 30000.0 51250.0
Time taken: 137.343 seconds, Fetched: 1 row(s)
hive>

```

- **Construction of view**

```

103 Goutam 80000.0
108 Sindhu 65000.0
Time taken: 148.852 seconds, Fetched: 2 row(s)
hive> select max(amount),min(amount),avg(amount);
FAILED: SemanticException [Error 10004]: Line 1:11 Invalid table alias or column reference 'amount': (possible column names are: )
hive> select max(amount) as MaxAnt,min(amount) as MinAnt,avg(amount) as AvgAnt;
FAILED: SemanticException [Error 10004]: Line 1:11 Invalid table alias or column reference 'amount': (possible column names are: )
hive> select max(amount) as MaxAnt,min(amount) as MinAnt,avg(amount) as AvgAnt from bank;
Query ID = hdoop_20210707133437_b7f61457-97ff-4825-aa2e-6a6e1f567806
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks determined at compile time: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1625638408330_0007, Tracking URL = http://asha-VirtualBox:8088/proxy/application_1625638408330_0007/
Kill Command = /home/hadoop/hadoop-3.2.1/bin/mapred job -kill job_1625638408330_0007
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2021-07-07 13:35:04,497 Stage-1 map = 0%, reduce = 0%
2021-07-07 13:36:04,657 Stage-1 map = 0%, reduce = 0%
2021-07-07 13:36:20,790 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 4.79 sec
2021-07-07 13:36:50,858 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 10.00 sec
MapReduce Total cumulative CPU time: 10 seconds 80 msec
Ended Job = job_1625638408330_0007
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 10.00 sec HDFS Read: 17237 HDFS Write: 123 SUCCESS
Total MapReduce CPU Time Spent: 10 seconds 80 msec
OK
80000.0 30000.0 51250.0
Time taken: 137.343 seconds, Fetched: 1 row(s)
hive> create view Cust_Transac as select tid,cname from bank;
OK
Time taken: 3.58 seconds
hive> select * from Cust_Transac;
OK
101 Asha
102 Ananya
103 Goutam
104 Nandita
105 Pavitra
106 Venugopal
107 Bharati
108 Sindhu
Time taken: 0.771 seconds, Fetched: 8 row(s)
hive>

```

- **Creating one more table called customer(cust\_id int,tid int,loc string)  
And importing values from csv file**

```

Time taken: 0.771 seconds, Fetched: 8 row(s)
hive> create table customer(cust_id int,tid int,loc string)row format delimited fields terminated by ",";
OK
Time taken: 0.647 seconds
hive> load data local inpath '/home/hadoop/cust.csv' into table customer;
Loading data to table customer.customer
OK
Time taken: 1.763 seconds
hive> select * from customer;
OK
201 102 Bangalore
202 105 Mysore
203 103 Dharwad
Time taken: 1.3 seconds, Fetched: 3 row(s)
hive>

```

- Experimenting on join operation

```

hadoop@asha-VirtualBox: ~/apache-hive-3.1.2-bin/conf
102 Ananya
103 Goutan
104 Nandita
105 Pavitra
106 Venugopal
107 Bharati
108 Sindhu
Time taken: 0.771 seconds, Fetched: 8 row(s)
hive> create table customer(cust_id int,tid int,loc string)row format delimited fields terminated by ",";
OK
Time taken: 0.647 seconds
hive> load data local inpath '/home/hadoop/cust.csv' into table customer;
Loading data to table customer.customer
OK
Time taken: 1.763 seconds
hive> select * from customer;
OK
201 102 Bangalore
202 105 Mysore
203 103 Dharwad
Time taken: 1.3 seconds, Fetched: 3 row(s)
hive> select cname,loc from bank_b,customer c where b.tid=c.tid;
Query ID = hadoop_20210707134344_8dee89f7-e82d-4cf2-ba92-cddbc032fec5
Total jobs = 1
SLF4J: Found binding in [jar:file:/home/hadoop/apache-hive-3.1.2-bin/lib/log4j-slf4j-impl-2.10.0.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/home/hadoop/hadoop-3.2.1/share/hadoop/common/lib/slf4j-log4j12-1.7.25.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
Execution completed successfully
MapReduce task succeeded
Launching Job 1 out of 1
Number of reduce tasks is set to 0 since there's no reduce operator
Starting Job = job_1625638408330_0008, Tracking URL = http://asha-VirtualBox:8088/proxy/application_1625638408330_0008/
Kill Command = /home/hadoop/hadoop-3.2.1/bin/mapred job -kill job_1625638408330_0008
Hadoop job information for Stage-3: number of mappers: 1; number of reducers: 0
2021-07-07 13:45:57,608 Stage-3 map = 0%, reduce = 0%
2021-07-07 13:46:28,693 Stage-3 map = 100%, reduce = 0%, Cumulative CPU 5.85 sec
MapReduce Total cumulative CPU time: 5 seconds 850 msec
Ended Job = job_1625638408330_0008
MapReduce Jobs Launched:
Stage-Stage-3: Map: 1 Cumulative CPU: 5.85 sec HDFS Read: 9511 HDFS Write: 170 SUCCESS
Total MapReduce CPU Time Spent: 5 seconds 850 msec
OK
Ananya Bangalore
Goutan Dharwad
Pavitra Mysore
Time taken: 169.712 seconds, Fetched: 3 row(s)
hive>
>

```

- Experimenting on alter operation

```

hadoop@asha-VirtualBox: ~/apache-hive-3.1.2-bin/conf
Time taken: 0.996 seconds, Fetched: 3 row(s)
hive> alter table cust add columns(phno int);
OK
Time taken: 1.786 seconds
hive> desc cust;
OK
cust_id      int
tid          int
loc          string
phno         int
Time taken: 0.697 seconds, Fetched: 4 row(s)
hive> insert into cust values(204,101,"Sirsi",9482813602);
Query ID = hadoop_20210707135108_e7ab72a7-c134-476d-97c3-724e2cbea440
Total jobs = 3
Launching Job 1 out of 3
Number of reduce tasks determined at compile time: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1625638408330_0009, Tracking URL = http://asha-VirtualBox:8088/proxy/application_1625638408330_0009/
Kill Command = /home/hadoop/hadoop-3.2.1/bin/mapred job -kill job_1625638408330_0009
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2021-07-07 13:51:56,535 Stage-1 map = 0%, reduce = 0%
2021-07-07 13:52:42,637 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 5.97 sec
2021-07-07 13:53:27,299 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 9.8 sec
MapReduce Total cumulative CPU time: 9 seconds 800 msec
Ended Job = job_1625638408330_0009
Stage-4 is selected by condition resolver.
Stage-3 is filtered out by condition resolver.
Stage-5 is filtered out by condition resolver.
Moving data to directory hdfs://127.0.0.1:9000/user/hive/warehouse/customer.db/cust/.hive-staging_hive_2021-07-07_13-51-08_667_8772405825846635967-1/-ext-10000
Loading data to table customer.cust
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 9.8 sec HDFS Read: 17846 HDFS Write: 341 SUCCESS
Total MapReduce CPU Time Spent: 9 seconds 800 msec
OK
Time taken: 144.235 seconds
hive> select * from cust;
OK
204 101 Sirsi 892879010
201 102 Bangalore NULL
202 105 Mysore NULL
203 103 Dharwad NULL
Time taken: 2.466 seconds, Fetched: 4 row(s)
hive>

```