

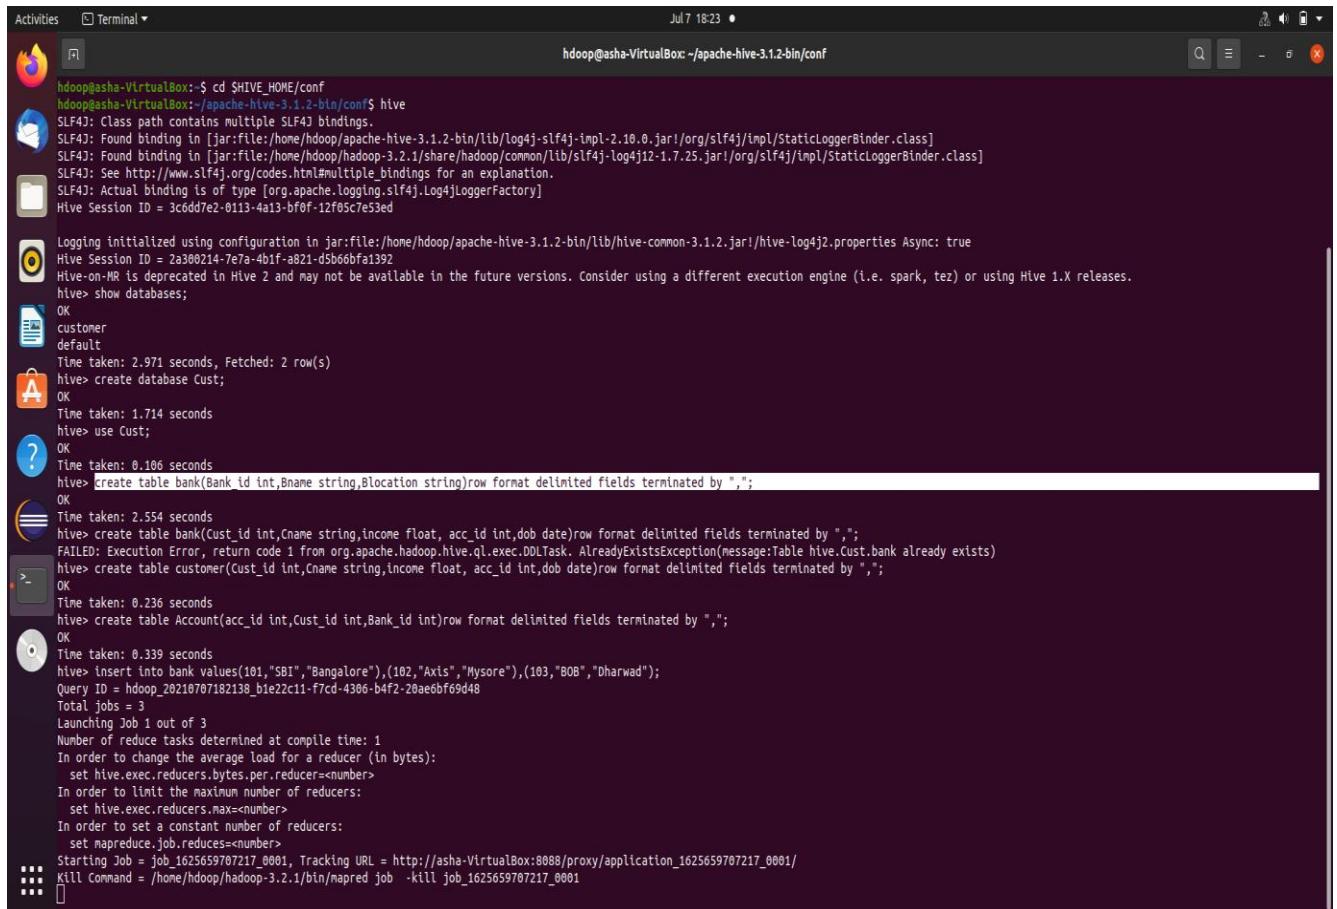
BIG DATA LAB HIVE ASSIGNMENT 2

Name : Asha Venugopal Madguni

USN : 1NT18IS033

Sem and Sec : 6-'A' (A2 Batch)

Create table Bank(Bank_id:integer,Bname:string,blocation:string)



The screenshot shows a terminal window titled "Terminal" with the command "hive" entered. The session starts with several SLF4J binding logs:

```
SLF4J: Class path contains multiple SLF4J bindings.  
SLF4J: Found binding in [jar:file:/home/hadoop/apache-hive-3.1.2-bin/lib/log4j-slf4j-impl-2.10.0.jar!/org/slf4j/impl/StaticLoggerBinder.class]  
SLF4J: Found binding in [jar:file:/home/hadoop/hadoop-3.2.1/share/hadoop/common/lib/slf4j-log4j12-1.7.25.jar!/org/slf4j/impl/StaticLoggerBinder.class]  
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.  
SLF4J: Actual binding is of type [org.apache.logging.slf4j.Log4jLoggerFactory]
```

Hive Session ID is set to 3c6dd7e2-0113-4a13-bff-12f05c7e53ed.

Logging is initialized using configuration in jar:file:/home/hadoop/apache-hive-3.1.2-bin/lib/hive-common-3.1.2.jar!/hive-log4j2.properties Async: true. Hive Session ID is 2a300214-7e7a-4b1f-a821-d5b66bfa1392. Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine (i.e. spark, tez) or using Hive 1.X releases.

hive> show databases;

OK

customer
default

Time taken: 2.971 seconds, Fetched: 2 row(s)

hive> create database Cust;

OK

Time taken: 1.714 seconds

hive> use Cust;

OK

Time taken: 0.106 seconds

hive> create table bank(Bank_id int,Bname string,Blocation string)row format delimited fields terminated by ",";

OK

Time taken: 2.554 seconds

hive> create table bank(Cust_id int,Cname string,income float, acc_id int,dob date)row format delimited fields terminated by ",";

FAILED: Execution Error, return code 1 from org.apache.hadoop.hive.qll.exec.DDLTask. AlreadyExistsException(message:Table hive.Cust.bank already exists)

hive> create table customer(Cust_id int,Cname string,income float, acc_id int,dob date)row format delimited fields terminated by ",";

OK

Time taken: 0.236 seconds

hive> create table Account(acc_id int,Cust_id int,Bank_id int)row format delimited fields terminated by ",";

OK

Time taken: 0.339 seconds

hive> insert into bank values(101,"SBI","Bangalore"),(102,"Axis","Mysore"),(103,"BOB","Dharwad");

Query ID = hdoop_20210707182138_b1e22c11-f7cd-4306-b4f2-20ae6bf69d48

Total jobs = 3

Launching Job 1 out of 3

Number of reduce tasks determined at compile time: 1

In order to change the average load for a reducer (in bytes):
set hive.exec.reducers.bytes.per.reducer=<number>

In order to limit the maximum number of reducers:
set hive.exec.reducers.max=<number>

In order to set a constant number of reducers:
set mapreduce.job.reduces=<number>

Starting Job = job_1625659707217_0001, Tracking URL = http://asha-VirtualBox:8088/proxy/application_1625659707217_0001/
Kill Command = /home/hadoop/hadoop-3.2.1/bin/mapred job -kill job_1625659707217_0001

Create table

customer(Cust_id:integer,Cname:string,income:float,acc_id:int,dob:date)

```
Activities Terminal Jul 7 18:23 • hdoop@asha-VirtualBox: ~/apache-hive-3.1.2-bin/conf
hdoop@asha-VirtualBox:~/apache-hive-3.1.2-bin/conf$ hive
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/home/hadoop/apache-hive-3.1.2-bin/lib/log4j-slf4j-impl-2.10.0.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/home/hadoop/hadoop-3.2.1/share/hadoop/common/lib/slf4j-log4j12-1.7.25.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.apache.logging.slf4j.Log4jLoggerFactory]
Hive Session ID = 3c6dd7e2-0113-4a13-bf0f-12f05c7e53ed

Logging initialized using configuration in jar:file:/home/hadoop/apache-hive-3.1.2-bin/lib/hive-common-3.1.2.jar!/hive-log4j2.properties Async: true
Hive Session ID = 2a300214-7e7a-4b1f-a821-d5b66fa1392
Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine (i.e. spark, tez) or using Hive 1.X releases.
hive> show databases;
OK
customer
default
Time taken: 2.971 seconds, Fetched: 2 row(s)
hive> create database Cust;
OK
Time taken: 1.714 seconds
hive> use Cust;
OK
Time taken: 0.106 seconds
hive> create table bank(Bank_id int,Bname string,Blocation string)row format delimited fields terminated by ",";
OK
Time taken: 2.554 seconds
hive> create table bank(Cust_id int,Cname string,income float, acc_id int,dob date)row format delimited fields terminated by ",";
FAILED: Execution Error, return code 1 from org.apache.hadoop.hive.ql.exec.DDLTask. AlreadyExistsException(message:Table hive.Cust.bank already exists)
hive> create table customer(Cust_id int,Cname string,income float, acc_id int,dob date)row format delimited fields terminated by ",";
OK
Time taken: 0.236 seconds
hive> create table Account(acc_id int,Cust_id int,Bank_id int)row format delimited fields terminated by ",";
OK
Time taken: 0.339 seconds
hive> insert into bank values(101,"SBI","Bangalore"),(102,"Axis","Mysore"),(103,"BOB","Dharwad");
Query ID = hdoop_20210707182138_b1e22c11-f7cd-4306-b4f2-20ae6bf69d48
Total jobs = 3
Launching Job 1 out of 3
Number of reduce tasks determined at compile time: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1625659707217_0001, Tracking URL = http://asha-VirtualBox:8088/proxy/application_1625659707217_0001
Kill Command = /home/hadoop/hadoop-3.2.1/bin/mapred job -kill job_1625659707217_0001
```

Create table Accont(acc_id:int,cust_id:int,bank_id:int)

```
Activities Terminal Jul 7 18:23 • hdoop@asha-VirtualBox: ~/apache-hive-3.1.2-bin/conf
hdoop@asha-VirtualBox:~/apache-hive-3.1.2-bin/conf$ hive
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/home/hadoop/apache-hive-3.1.2-bin/lib/log4j-slf4j-impl-2.10.0.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/home/hadoop/hadoop-3.2.1/share/hadoop/common/lib/slf4j-log4j12-1.7.25.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.apache.logging.slf4j.Log4jLoggerFactory]
Hive Session ID = 3c6dd7e2-0113-4a13-bf0f-12f05c7e53ed

Logging initialized using configuration in jar:file:/home/hadoop/apache-hive-3.1.2-bin/lib/hive-common-3.1.2.jar!/hive-log4j2.properties Async: true
Hive Session ID = 2a300214-7e7a-4b1f-a821-d5b66fa1392
Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine (i.e. spark, tez) or using Hive 1.X releases.
hive> show databases;
OK
customer
default
Time taken: 2.971 seconds, Fetched: 2 row(s)
hive> create database Cust;
OK
Time taken: 1.714 seconds
hive> use Cust;
OK
Time taken: 0.106 seconds
hive> create table bank(Bank_id int,Bname string,Blocation string)row format delimited fields terminated by ",";
OK
Time taken: 2.554 seconds
hive> create table bank(Cust_id int,Cname string,income float, acc_id int,dob date)row format delimited fields terminated by ",";
FAILED: Execution Error, return code 1 from org.apache.hadoop.hive.ql.exec.DDLTask. AlreadyExistsException(message:Table hive.Cust.bank already exists)
hive> create table customer(Cust_id int,Cname string,income float, acc_id int,dob date)row format delimited fields terminated by ",";
OK
Time taken: 0.236 seconds
hive> create table Account(acc_id int,Cust_id int,Bank_id int)row format delimited fields terminated by ",";
OK
Time taken: 0.339 seconds
hive> insert into bank values(101,"SBI","Bangalore"),(102,"Axis","Mysore"),(103,"BOB","Dharwad");
Query ID = hdoop_20210707182138_b1e22c11-f7cd-4306-b4f2-20ae6bf69d48
Total jobs = 3
Launching Job 1 out of 3
Number of reduce tasks determined at compile time: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1625659707217_0001, Tracking URL = http://asha-VirtualBox:8088/proxy/application_1625659707217_0001
Kill Command = /home/hadoop/hadoop-3.2.1/bin/mapred job -kill job_1625659707217_0001
```

Insert 5 records using INSERT command

Inserting into bank table

```
Activities Terminal ▾ hdoop@asha-VirtualBox: ~/apache-hive-3.1.2-bin/conf
OK
Time taken: 1.714 seconds
hive> use Cust;
OK
Time taken: 0.106 seconds
hive> create table bank(Bank_id int,Bname string,Blocation string)row format delimited fields terminated by ",";
OK
Time taken: 2.554 seconds
hive> create table bank(Cust_id int,Cname string,income float, acc_id int,dob date)row format delimited fields terminated by ",";
FAILED: Execution Error, return code 1 from org.apache.hadoop.hive.ql.exec.DDLTask. AlreadyExistsException(message:Table hive.Cust.bank already exists)
hive> create table customer(Cust_id int,Cname string,income float, acc_id int,dob date)row format delimited fields terminated by ",";
OK
Time taken: 0.236 seconds
hive> create table Account(acc_id int,Cust_id int,Bank_id int)row format delimited fields terminated by ",";
OK
Time taken: 0.339 seconds
hive> insert into bank values(101,"SBI","Bangalore"),(102,"Axis","Mysore"),(103,"BOB","Dharwad");
Query ID = hdoop_20210707182138_bie22c11-f7cd-4306-b4f2-20ae6bf69d48
Total jobs = 3
Launching Job 1 out of 3
Number of reduce tasks determined at compile time: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1625659707217_0001, Tracking URL = http://asha-VirtualBox:8088/proxy/application_1625659707217_0001/
Kill Command = /home/hdoop/hadoop-3.2.1/bin/mapred job -kill job_1625659707217_0001
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2021-07-07 18:23:43,467 Stage-1 map = 0%, reduce = 0%
2021-07-07 18:24:22,726 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 4.54 sec
2021-07-07 18:24:50,571 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 7.85 sec
MapReduce Total cumulative CPU time: 7 seconds 850 msec
Ended Job = job_1625659707217_0001
Stage-4 is selected by condition resolver.
Stage-3 is filtered out by condition resolver.
Stage-5 is filtered out by condition resolver.
Moving data to directory hdfs://127.0.0.1:9000/user/hive/warehouse/cust.db/bank/.hive-staging_hive_2021-07-07_18-21-38_369_3597759571545103119-1/-ext-10000
Loading data to table cust.bank
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 7.85 sec HDFS Read: 16904 HDFS Write: 382 SUCCESS
Total MapReduce CPU Time Spent: 7 seconds 850 msec
OK
Time taken: 198.6 seconds
hive> insert into bank values(104,"SBI","Yellapur"),(105,"Axis","Sirsi");
Query ID = hdoop_20210707182652_e2dc077e-bc0a-4ed6-ac5-4f3c4afa76b4
Total jobs = 3
Launching Job 1 out of 3
```

```
Activities Terminal ▾ hdoop@asha-VirtualBox: ~/apache-hive-3.1.2-bin/conf
Stage-4 is selected by condition resolver.
Stage-3 is filtered out by condition resolver.
Stage-5 is filtered out by condition resolver.
Moving data to directory hdfs://127.0.0.1:9000/user/hive/warehouse/cust.db/bank/.hive-staging_hive_2021-07-07_18-21-38_369_3597759571545103119-1/-ext-10000
Loading data to table cust.bank
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 7.85 sec HDFS Read: 16904 HDFS Write: 382 SUCCESS
Total MapReduce CPU Time Spent: 7 seconds 850 msec
OK
Time taken: 198.6 seconds
hive> insert into bank values(104,"SBI","Yellapur"),(105,"Axis","Sirsi");
Query ID = hdoop_20210707182652_e2dc077e-bc0a-4ed6-ac5-4f3c4afa76b4
Total jobs = 3
Launching Job 1 out of 3
Number of reduce tasks determined at compile time: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1625659707217_0002, Tracking URL = http://asha-VirtualBox:8088/proxy/application_1625659707217_0002/
Kill Command = /home/hdoop/hadoop-3.2.1/bin/mapred job -kill job_1625659707217_0002
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2021-07-07 18:27:24,181 Stage-1 map = 0%, reduce = 0%
2021-07-07 18:27:45,241 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 4.39 sec
2021-07-07 18:28:14,932 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 7.8 sec
MapReduce Total cumulative CPU time: 7 seconds 800 msec
Ended Job = job_1625659707217_0002
Stage-4 is selected by condition resolver.
Stage-3 is filtered out by condition resolver.
Stage-5 is filtered out by condition resolver.
Moving data to directory hdfs://127.0.0.1:9000/user/hive/warehouse/cust.db/bank/.hive-staging_hive_2021-07-07_18-26-52_715_7976836924505372864-1/-ext-10000
Loading data to table cust.bank
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 7.8 sec HDFS Read: 16809 HDFS Write: 322 SUCCESS
Total MapReduce CPU Time Spent: 7 seconds 800 msec
OK
Time taken: 85.843 seconds
hive> select * from bank;
OK
101    SBI    Bangalore
102    Axis   Mysore
103    BOB    Dharwad
104    SBI    Yellapur
105    Axis   Sirsi
Time taken: 1.763 seconds, Fetched: 5 row(s)
hive> [ ]
```

Inserting into customer table

```
Activities Terminal •
hadoop@asha-VirtualBox: ~/apache-hive-3.1.2-bin/conf
Loading data to table cust.bank
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 7.8 sec HDFS Read: 16809 HDFS Write: 322 SUCCESS
Total MapReduce CPU Time Spent: 7 seconds 800 msec
OK
Time taken: 85.843 seconds
hive> select * from bank;
OK
101 SBI Bangalore
102 Axis Mysore
103 BOB Dharwad
104 SBI Yellapur
105 Axis Sirsi
Time taken: 1.763 seconds. Fetched: 5 row(s)
hive> insert into customer values(201,"Asha",30000.00,301,2000-07-27),(202,"Goutam",80000.00,302,2000-11-16),(203,"Bharati",50000.00,303,1991-06-26);
FAILED: UDFArgumentException CAST as DATE only allows date,string, or timestamp types
hive> insert into customer values(201,"Asha",30000.00,301),(202,"Goutam",80000.00,302,"2000-07-27"),(203,"Bharati",50000.00,303,"1991-06-26");
Query ID = hdoop_20210707183511_98bd7abd-63f4-461d-b9ff-33f6fbad51a
Total jobs = 3
Launching Job 1 out of 3
Number of reduce tasks is set to 0 since there's no reduce operator
Starting Job = job_1625659707217_0003, Tracking URL = http://asha-VirtualBox:8088/proxy/application_1625659707217_0003/
Kill Command = /home/hadoop/hadoop-3.2.1/bin/mapred job -kill job_1625659707217_0003
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 0
2021-07-07 18:35:40,187 Stage-1 map = 0%, reduce = 0%
2021-07-07 18:35:58,528 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 4.2 sec
MapReduce Total cumulative CPU time: 4 seconds 200 msec
Ended Job = job_1625659707217_0003
Stage-4 is selected by condition resolver.
Stage-3 is filtered out by condition resolver.
Stage-5 is filtered out by condition resolver.
Moving data to directory hdfs://127.0.0.1:9000/user/hive/warehouse/cust.db/customer/.hive-staging_hive_2021-07-07_18-35-11_617_8710127301975818774-1/-ext-10000
Loading data to table cust.customer
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Cumulative CPU: 4.2 sec HDFS Read: 6686 HDFS Write: 170 SUCCESS
Total MapReduce CPU Time Spent: 4 seconds 200 msec
OK
Time taken: 49.918 seconds
hive> insert into customer values(204,"Pavitra",60000.00,304,"1990-07-26"),(205,"Nandita",70000.00,305,"1992-06-15");
Query ID = hdoop_20210707183728_f6a70863-391f-4818-9788-04c6981f1902
Total jobs = 3
Launching Job 1 out of 3
Number of reduce tasks is set to 0 since there's no reduce operator
Starting Job = job_1625659707217_0004, Tracking URL = http://asha-VirtualBox:8088/proxy/application_1625659707217_0004/
Kill Command = /home/hadoop/hadoop-3.2.1/bin/mapred job -kill job_1625659707217_0004
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 0
2021-07-07 18:35:40,187 Stage-1 map = 0%, reduce = 0%
2021-07-07 18:35:58,528 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 4.2 sec
MapReduce Total cumulative CPU time: 4 seconds 200 msec
Ended Job = job_1625659707217_0004
Stage-4 is selected by condition resolver.
Stage-3 is filtered out by condition resolver.
Stage-5 is filtered out by condition resolver.
Moving data to directory hdfs://127.0.0.1:9000/user/hive/warehouse/cust.db/customer/.hive-staging_hive_2021-07-07_18-35-11_617_8710127301975818774-1/-ext-10000
Loading data to table cust.customer
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Cumulative CPU: 4.2 sec HDFS Read: 6686 HDFS Write: 170 SUCCESS
Total MapReduce CPU Time Spent: 4 seconds 200 msec
OK
Time taken: 49.918 seconds
hive>
```

```
Activities Terminal •
hadoop@asha-VirtualBox: ~/apache-hive-3.1.2-bin/conf
Starting Job = job_1625659707217_0003, Tracking URL = http://asha-VirtualBox:8088/proxy/application_1625659707217_0003/
Kill Command = /home/hadoop/hadoop-3.2.1/bin/mapred job -kill job_1625659707217_0003
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 0
2021-07-07 18:35:40,187 Stage-1 map = 0%, reduce = 0%
2021-07-07 18:35:58,528 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 4.2 sec
MapReduce Total cumulative CPU time: 4 seconds 200 msec
Ended Job = job_1625659707217_0003
Stage-4 is selected by condition resolver.
Stage-3 is filtered out by condition resolver.
Stage-5 is filtered out by condition resolver.
Moving data to directory hdfs://127.0.0.1:9000/user/hive/warehouse/cust.db/customer/.hive-staging_hive_2021-07-07_18-35-11_617_8710127301975818774-1/-ext-10000
Loading data to table cust.customer
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Cumulative CPU: 4.2 sec HDFS Read: 6686 HDFS Write: 170 SUCCESS
Total MapReduce CPU Time Spent: 4 seconds 200 msec
OK
Time taken: 49.918 seconds
hive> insert into customer values(204,"Pavitra",60000.00,304,"1990-07-26"),(205,"Nandita",70000.00,305,"1992-06-15");
Query ID = hdoop_20210707183728_f6a70863-391f-4818-9788-04c6981f1902
Total jobs = 3
Launching Job 1 out of 3
Number of reduce tasks is set to 0 since there's no reduce operator
Starting Job = job_1625659707217_0004, Tracking URL = http://asha-VirtualBox:8088/proxy/application_1625659707217_0004/
Kill Command = /home/hadoop/hadoop-3.2.1/bin/mapred job -kill job_1625659707217_0004
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 0
2021-07-07 18:37:53,422 Stage-1 map = 0%, reduce = 0%
2021-07-07 18:38:10,654 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 4.44 sec
MapReduce Total cumulative CPU time: 4 seconds 440 msec
Ended Job = job_1625659707217_0004
Stage-4 is selected by condition resolver.
Stage-3 is filtered out by condition resolver.
Stage-5 is filtered out by condition resolver.
Moving data to directory hdfs://127.0.0.1:9000/user/hive/warehouse/cust.db/customer/.hive-staging_hive_2021-07-07_18-37-28_780_4605519924051754910-1/-ext-10000
Loading data to table cust.customer
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Cumulative CPU: 4.44 sec HDFS Read: 6508 HDFS Write: 139 SUCCESS
Total MapReduce CPU Time Spent: 4 seconds 440 msec
OK
Time taken: 45.215 seconds
hive> select * from customer;
OK
201 Asha 30000.0 301 2000-07-27
202 Goutam 80000.0 302 2000-11-16
203 Bharati 50000.0 303 1991-06-26
204 Pavitra 60000.0 304 1990-07-26
205 Nandita 70000.0 305 1992-06-15
Time taken: 0.589 seconds, Fetched: 5 row(s)
hive> 
```

Inserting into Account table

```
Activities Terminal Jul 7 18:50 • hdoop@asha-VirtualBox: ~/apache-hive-3.1.2-bin/conf
at org.apache.hadoop.hive.cli.CliDriver.executeDriver(CliDriver.java:821)
at org.apache.hadoop.hive.cli.CliDriver.run(CliDriver.java:759)
at org.apache.hadoop.hive.cli.CliDriver.main(CliDriver.java:683)
at sun.reflect.NativeMethodAccessorImpl.invoke0(Native Method)
at sun.reflect.NativeMethodAccessorImpl.invoke(NativeMethodAccessorImpl.java:62)
at sun.reflect.DelegatingMethodAccessorImpl.invoke(DelegatingMethodAccessorImpl.java:43)
at java.lang.reflect.Method.invoke(Method.java:498)
at org.apache.hadoop.util.RunJar.run(RunJar.java:323)
at org.apache.hadoop.util.RunJar.main(RunJar.java:236)
FAILED: ParseException line 1:28 cannot recognize input near 'Account' '[' 'where' in table name
hive> delete from Account WHERE acc_id=101;
FAILED: SemanticException [Error 10294]: Attempt to do update or delete using transaction manager that does not support these operations.
hive> DELETE from Account WHERE acc_id=101;
FAILED: SemanticException [Error 10294]: Attempt to do update or delete using transaction manager that does not support these operations.
hive> drop table Account;
OK
Time taken: 2.074 seconds
hive> show tables;
OK
bank
customer
Time taken: 0.226 seconds, Fetched: 2 row(s)
hive> create table Account(acc_id int,Cust_id int,Bank_id int)row format delimited fields terminated by ",";
OK
Time taken: 0.245 seconds
hive> show tables;
OK
account
bank
customer
Time taken: 0.141 seconds, Fetched: 3 row(s)
hive> insert into Account values(301,202,101),(302,203,102),(303,201,104);
Query ID = hdoop_20210707184933_552839b4-ab38-4f5d-b60c-2de0f26d275b
Total jobs = 3
Launching Job 1 out of 3
Number of reduce tasks determined at compile time: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1625659707217_0006, Tracking URL = http://asha-VirtualBox:8088/proxy/application_1625659707217_0006/
Kill Command = /home/hdoop/hadoop-3.2.1/bin/mapred job -kill job_1625659707217_0006
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2021-07-07 18:49:57,672 Stage-1 map = 0%,  reduce = 0%
2021-07-07 18:50:14,133 Stage-1 map = 100%,  reduce = 0%, Cumulative CPU 4.5 sec
[...]
```

```
Activities Terminal Jul 7 18:52 • hdoop@asha-VirtualBox: ~/apache-hive-3.1.2-bin/conf
Stage-4 is selected by condition resolver.
Stage-3 is filtered out by condition resolver.
Stage-5 is filtered out by condition resolver.
Moving data to directory hdfs://127.0.0.1:9000/user/hive/warehouse/cust.db/account/.hive-staging_hive_2021-07-07_18-49-33_317_1598376274790812557-1/-ext-10000
Loading data to table cust.account
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1  Cumulative CPU: 7.87 sec  HDFS Read: 15710 HDFS Write: 342 SUCCESS
Total MapReduce CPU Time Spent: 7 seconds 870 msec
OK
Time taken: 67.545 seconds
hive> insert into Account values(304,202,103),(305,203,105);
Query ID = hdoop_20210707185122_c75489a5-d417-4d07-9804-884b8dfe0dc3
Total jobs = 3
Launching Job 1 out of 3
Number of reduce tasks determined at compile time: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1625659707217_0007, Tracking URL = http://asha-VirtualBox:8088/proxy/application_1625659707217_0007/
Kill Command = /home/hdoop/hadoop-3.2.1/bin/mapred job -kill job_1625659707217_0007
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2021-07-07 18:51:43,953 Stage-1 map = 0%,  reduce = 0%
2021-07-07 18:52:02,706 Stage-1 map = 100%,  reduce = 0%, Cumulative CPU 4.24 sec
2021-07-07 18:52:17,688 Stage-1 map = 100%,  reduce = 100%, Cumulative CPU 7.65 sec
MapReduce Total cumulative CPU time: 7 seconds 650 msec
Ended Job = job_1625659707217_0007
Stage-4 is selected by condition resolver.
Stage-3 is filtered out by condition resolver.
Stage-5 is filtered out by condition resolver.
Moving data to directory hdfs://127.0.0.1:9000/user/hive/warehouse/cust.db/account/.hive-staging_hive_2021-07-07_18-51-22_559_3622744112434897685-1/-ext-10000
Loading data to table cust.account
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1  Cumulative CPU: 7.65 sec  HDFS Read: 15641 HDFS Write: 313 SUCCESS
Total MapReduce CPU Time Spent: 7 seconds 650 msec
OK
Time taken: 58.13 seconds
hive> select * from Account;
OK
301 202 101
302 203 102
303 201 104
304 202 103
305 203 105
Time taken: 0.47 seconds, Fetched: 5 row(s)
hive> [...]
```

Rename the table Account to Accounts

```
Activities Terminal • hdoop@asha-VirtualBox: ~/apache-hive-3.1.2-bin/conf
Time taken: 67.545 seconds
hive> insert into Account values(304,202,103),(305,203,105);
Query ID = hdoop_20210707185122_c75489a5-d417-4d07-9804-884b8dfe0dc3
Total jobs = 3
Launching Job 1 out of 3
Number of reduce tasks determined at compile time: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1625659707217_0007, Tracking URL = http://asha-VirtualBox:8088/proxy/application_1625659707217_0007/
Kill Command = /home/hdoop/hadoop-3.2.1/bin/mapred job -kill job_1625659707217_0007
2021-07-07 18:51:43,953 Stage-1 map = 0%, reduce = 0%
2021-07-07 18:52:02,768 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 4.24 sec
2021-07-07 18:52:17,681 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 7.65 sec
MapReduce Total cumulative CPU time: 7 seconds 650 msec
Ended Job = job_1625659707217_0007
Stage-4 is selected by condition resolver.
Stage-3 is filtered out by condition resolver.
Stage-5 is filtered out by condition resolver.
Moving data to directory hdfs://127.0.0.1:9000/user/hive/warehouse/cust.db/account/.hive-staging_hive_2021-07-07_18-51-22_559_3622744112434897685-1/-ext-10000
Loading data to table cust.account
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 7.65 sec HDFS Read: 15641 HDFS Write: 313 SUCCESS
Total MapReduce CPU Time Spent: 7 seconds 650 msec
OK
Time taken: 58.13 seconds
hive> select * from Account;
OK
301 202 101
302 203 102
303 201 104
304 202 103
305 203 105
Time taken: 0.47 seconds, Fetched: 5 row(s)
hive> alter table Account rename to Accounts;
OK
Time taken: 0.529 seconds
hive> show tables;
OK
accounts
bank
customer
Time taken: 0.135 seconds, Fetched: 3 row(s)
hive> 
```

Rename the Column “blocation” to “location”

```
Activities Terminal • hdoop@asha-VirtualBox: ~/apache-hive-3.1.2-bin/conf
Time taken: 67.545 seconds
hive> set hive.exec.reducers.max=<number>;
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1625659707217_0007, Tracking URL = http://asha-VirtualBox:8088/proxy/application_1625659707217_0007/
Kill Command = /home/hdoop/hadoop-3.2.1/bin/mapred job -kill job_1625659707217_0007
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2021-07-07 18:51:43,953 Stage-1 map = 0%, reduce = 0%
2021-07-07 18:52:02,768 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 4.24 sec
2021-07-07 18:52:17,681 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 7.65 sec
MapReduce Total cumulative CPU time: 7 seconds 650 msec
Ended Job = job_1625659707217_0007
Stage-4 is selected by condition resolver.
Stage-3 is filtered out by condition resolver.
Stage-5 is filtered out by condition resolver.
Moving data to directory hdfs://127.0.0.1:9000/user/hive/warehouse/cust.db/account/.hive-staging_hive_2021-07-07_18-51-22_559_3622744112434897685-1/-ext-10000
Loading data to table cust.account
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 7.65 sec HDFS Read: 15641 HDFS Write: 313 SUCCESS
Total MapReduce CPU Time Spent: 7 seconds 650 msec
OK
Time taken: 58.13 seconds
hive> select * from Account;
OK
301 202 101
302 203 102
303 201 104
304 202 103
305 203 105
Time taken: 0.47 seconds, Fetched: 5 row(s)
hive> alter table Account rename to Accounts;
OK
Time taken: 0.529 seconds
hive> show tables;
OK
accounts
bank
customer
Time taken: 0.135 seconds, Fetched: 3 row(s)
hive> alter table bank change blocation location string;
OK
Time taken: 0.348 seconds
hive> desc bank;
OK
bank_id      int
bname        string
location      string
Time taken: 0.147 seconds, Fetched: 3 row(s)
hive> 
```

Retrieve all the customers who have account in Axis bank

```
Activities Terminal Jul 7 19:02 • hdoop@asha-VirtualBox: ~/apache-hive-3.1.2-bin/conf

Total jobs = 1
Execution completed successfully
MapredLocal task succeeded
Launching Job 1 out of 1
Number of reduce tasks is set to 0 since there's no reduce operator
Starting Job = job_1625659707217_0008, Tracking URL = http://asha-VirtualBox:8088/proxy/application_1625659707217_0008/
Kill Command = /home/hdoop/hadoop-3.2.1/bin/mapred job -kill job_1625659707217_0008
Hadoop job information for Stage-0: number of Mappers: 1; number of reducers: 0
2021-07-07 18:59:29.919 Stage-0 map = 0%, reduce = 0%
2021-07-07 18:59:47.457 Stage-0 map = 100%, reduce = 0%, Cumulative CPU 4.57 sec
MapReduce Total cumulative CPU time: 4 seconds 570 msec
Ended Job = job_1625659707217_0008
MapReduce Jobs Launched:
Stage-Stage-0: Map: 1 Cumulative CPU: 4.57 sec HDFS Read: 11921 HDFS Write: 135 SUCCESS
Total MapReduce CPU Time Spent: 4 seconds 570 msec
OK
203 Bharati
203 Bharati
Time taken: 81.199 seconds, Fetched: 2 row(s)
hive> select distinct c.Cust_Id,c.cname from customer c,bank b,accounts a where b.bank_id=a.bank_id and c.cust_id=a.cust_id and b.bnname="Axis";
No Stats for customer, Columns: cname, cust_id
Query ID = hdoop_20210707190029_0ec2d682-a7c-4207-9036-1de1ee93ba84
Total jobs = 1
Execution completed successfully
MapredLocal task succeeded
Launching Job 1 out of 1
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1625659707217_0009, Tracking URL = http://asha-VirtualBox:8088/proxy/application_1625659707217_0009/
Kill Command = /home/hdoop/hadoop-3.2.1/bin/mapred job -kill job_1625659707217_0009
Hadoop job information for Stage-0: number of mappers: 1; number of reducers: 1
2021-07-07 19:01:19.468 Stage-0 map = 0%, reduce = 0%
2021-07-07 19:01:35.728 Stage-0 map = 100%, reduce = 0%, Cumulative CPU 4.38 sec
2021-07-07 19:01:51.843 Stage-0 map = 100%, reduce = 100%, Cumulative CPU 8.45 sec
MapReduce Total cumulative CPU time: 8 seconds 450 msec
Ended Job = job_1625659707217_0009
MapReduce Jobs Launched:
Stage-Stage-0: Map: 1 Reduce: 1 Cumulative CPU: 8.45 sec HDFS Read: 21940 HDFS Write: 111 SUCCESS
Total MapReduce CPU Time Spent: 8 seconds 450 msec
OK
203 Bharati
Time taken: 84.39 seconds, Fetched: 1 row(s)
hive>
```

Retrieve the customer with maximum and minimum income

```
Activities Terminal Jul 7 19:35 • hdoop@asha-VirtualBox: ~/apache-hive-3.1.2-bin/conf

at org.apache.hadoop.hive.ql.exec.mr.ExecReducer.reduce(ExecReducer.java:234)
... 7 more
Caused by: java.lang.ClassCastException: java.lang.Float cannot be cast to java.lang.Boolean
at org.apache.hadoop.hive.ql.exec.FilterOperator.process(FilterOperator.java:124)
at org.apache.hadoop.hive.ql.exec.Operator.baseForward(Operator.java:995)
at org.apache.hadoop.hive.ql.exec.Operator.forward(Operator.java:941)
at org.apache.hadoop.hive.ql.exec.Operator.forward(Operator.java:928)
at org.apache.hadoop.hive.ql.exec.GroupByOperator.forward(GroupByOperator.java:1050)
at org.apache.hadoop.hive.ql.exec.GroupByOperator.processKey(GroupByOperator.java:850)
at org.apache.hadoop.hive.ql.exec.GroupByOperator.processKey(GroupByOperator.java:724)
at org.apache.hadoop.hive.ql.exec.GroupByOperator.process(GroupByOperator.java:790)
... 8 more

FAILED: Execution Error, return code 2 from org.apache.hadoop.hive.ql.exec.mr.MapRedTask
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 3.55 sec HDFS Read: 7110 HDFS Write: 0 FAIL
Total MapReduce CPU Time Spent: 3 seconds 550 msec
hive> select cust_id,cname,income from customer group by cust_id having max(income);
FAILED: SemanticException [Error 10025]: Line 1:15 Expression not in GROUP BY key 'cname'
hive> select cust_id,cname,income from customer group by cust_id,cname having max(income);
FAILED: SemanticException [Error 10025]: Line 1:21 Expression not in GROUP BY key 'income'
hive> select max(income),min(income) from customer;
Query ID = hdoop_20210707193343_26db9c79-5f73-45b3-aeff-ba29babcafe8
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks determined at compile time: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1625659707217_0013, Tracking URL = http://asha-VirtualBox:8088/proxy/application_1625659707217_0013/
Kill Command = /home/hdoop/hadoop-3.2.1/bin/mapred job -kill job_1625659707217_0013
Hadoop job information for Stage-0: number of mappers: 1; number of reducers: 1
2021-07-07 19:34:10.554 Stage-0 map = 0%, reduce = 0%
2021-07-07 19:34:35.114 Stage-0 map = 100%, reduce = 0%, Cumulative CPU 3.73 sec
2021-07-07 19:34:56.125 Stage-0 map = 100%, reduce = 100%, Cumulative CPU 7.27 sec
MapReduce Total cumulative CPU time: 7 seconds 270 msec
Ended Job = job_1625659707217_0013
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 7.27 sec HDFS Read: 14424 HDFS Write: 115 SUCCESS
Total MapReduce CPU Time Spent: 7 seconds 270 msec
OK
80000.0 30000.0
Time taken: 74.951 seconds, Fetched: 1 row(s)
hive>
```

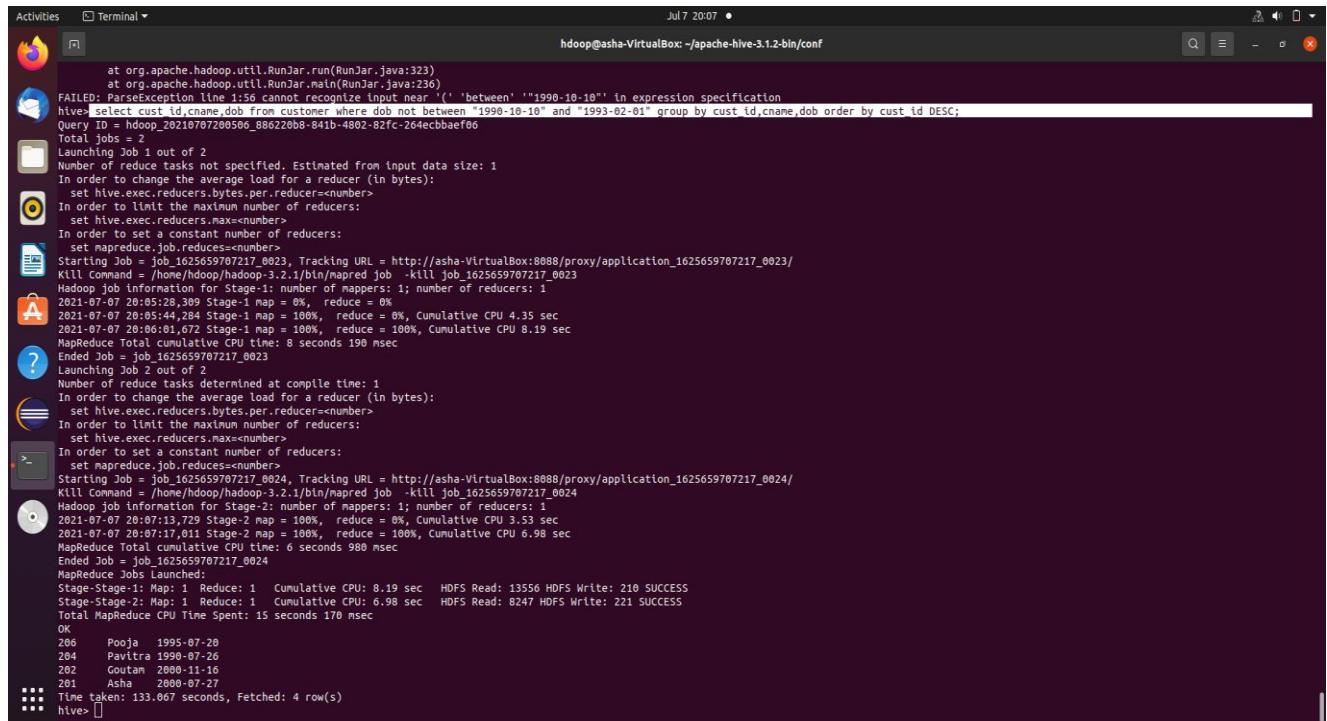
Retrieve all the customer who have account

```
Activities Terminal ▾
hadoop@asha-VirtualBox: ~/apache-hive-3.1.2-bin/conf
Number of reduce tasks is set to 0 since there's no reduce operator
Starting Job = job_1625659707217_0014, Tracking URL = http://asha-VirtualBox:8088/proxy/application_1625659707217_0014/
Kill Command = /home/hadoop/hadoop-3.2.1/bin/mapred job -kill job_1625659707217_0014
Hadoop job information for Stage-3: number of mappers: 1; number of reducers: 0
2021-07-07 19:37:47,229 Stage-3 map = 0%, reduce = 0%
2021-07-07 19:38:05,529 Stage-3 map = 100%, reduce = 0%, Cumulative CPU 4.46 sec
MapReduce Total cumulative CPU time: 4 seconds 460 msec
Ended Job = job_1625659707217_0014
MapReduce Jobs Launched:
Stage-Stage-3: Map: 1 Cumulative CPU: 4.46 sec HDFS Read: 9747 HDFS Write: 202 SUCCESS
Total MapReduce CPU Time Spent: 4 seconds 460 msec
OK
201 Asha
202 Goutam
202 Goutam
203 Bharati
203 Bharati
Time taken: 75.583 seconds, Fetched: 5 row(s)
hive> select distinct c.cust_id,c.cname from customer c,accounts a where c.cust_id=a.cust_id;
Query ID = hdoop_20210707193831_10007965-9eec-4d03-b6af-3f0930c8ead4
Total jobs = 1
Execution completed successfully
MapredLocal task succeeded
Launching Job 1 out of 1
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1625659707217_0015, Tracking URL = http://asha-VirtualBox:8088/proxy/application_1625659707217_0015/
Kill Command = /home/hadoop/hadoop-3.2.1/bin/mapred job -kill job_1625659707217_0015
Hadoop job information for Stage-2: number of mappers: 1; number of reducers: 1
2021-07-07 19:39:24,088 Stage-2 map = 0%, reduce = 0%
2021-07-07 19:39:43,218 Stage-2 map = 100%, reduce = 0%, Cumulative CPU 4.77 sec
2021-07-07 19:40:03,338 Stage-2 map = 100%, reduce = 100%, Cumulative CPU 8.48 sec
MapReduce Total cumulative CPU time: 8 seconds 480 msec
Ended Job = job_1625659707217_0015
MapReduce Jobs Launched:
Stage-Stage-2: Map: 1 Reduce: 1 Cumulative CPU: 8.48 sec HDFS Read: 18853 HDFS Write: 155 SUCCESS
Total MapReduce CPU Time Spent: 8 seconds 480 msec
OK
201 Asha
202 Goutam
203 Bharati
Time taken: 93.773 seconds, Fetched: 3 row(s)
hive>
```

Retrieve all the customers whose income is not more than 20000

```
Activities Terminal ▾
hadoop@asha-VirtualBox: ~/apache-hive-3.1.2-bin/conf
set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1625659707217_0015, Tracking URL = http://asha-VirtualBox:8088/proxy/application_1625659707217_0015/
Kill Command = /home/hadoop/hadoop-3.2.1/bin/mapred job -kill job_1625659707217_0015
Hadoop job information for Stage-2: number of mappers: 1; number of reducers: 1
2021-07-07 19:39:24,088 Stage-2 map = 0%, reduce = 0%
2021-07-07 19:39:43,218 Stage-2 map = 100%, reduce = 0%, Cumulative CPU 4.77 sec
2021-07-07 19:40:03,338 Stage-2 map = 100%, reduce = 100%, Cumulative CPU 8.48 sec
MapReduce Total cumulative CPU time: 8 seconds 480 msec
Ended Job = job_1625659707217_0015
Mapreduce Jobs Launched:
Stage-Stage-2: Map: 1 Reduce: 1 Cumulative CPU: 8.48 sec HDFS Read: 18853 HDFS Write: 155 SUCCESS
Total MapReduce CPU Time Spent: 8 seconds 480 msec
OK
201 Asha
202 Goutam
203 Bharati
Time taken: 93.773 seconds, Fetched: 3 row(s)
hive> Insert into customer values(206,"Pooja",10000.00,304,"1995-07-20");
Query ID = hdoop_20210707194117_938498a1-f659-4841-9ebcebc6069
Total jobs = 3
Launching Job 1 out of 3
Number of reduce tasks is set to 0 since there's no reduce operator
Starting Job = job_1625659707217_0016, Tracking URL = http://asha-VirtualBox:8088/proxy/application_1625659707217_0016/
Kill Command = /home/hadoop/hadoop-3.2.1/bin/mapred job -kill job_1625659707217_0016
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 0
2021-07-07 19:41:41,772 Stage-1 map = 0%, reduce = 0%
2021-07-07 19:41:59,978 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 4.18 sec
MapReduce Total cumulative CPU time: 4 seconds 180 msec
Ended Job = job_1625659707217_0016
Stage-4 is selected by condition resolver.
Stage-3 is filtered out by condition resolver.
Stage-5 is filtered out by condition resolver.
Moving data to directory hdfs://127.0.0.1:9000/user/hive/warehouse/cust.db/customer/.hive-staging_hive_2021-07-07_19-41-17_807_5514036148250295228-1/-ext-10000
Loading data to table cust.customer
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Cumulative CPU: 4.18 sec HDFS Read: 6396 HDFS Write: 102 SUCCESS
Total MapReduce CPU Time Spent: 4 seconds 180 msec
OK
Time taken: 46.626 seconds
hive> select cust_id,cname,income from customer where income<20000;
OK
206 Pooja 10000.0
Time taken: 1.239 seconds, Fetched: 1 row(s)
hive>
```

Retrieve all the customers whose dob not between 10-10-1990 and 1-2-1993 group based on cust_id and sort the results in descending order

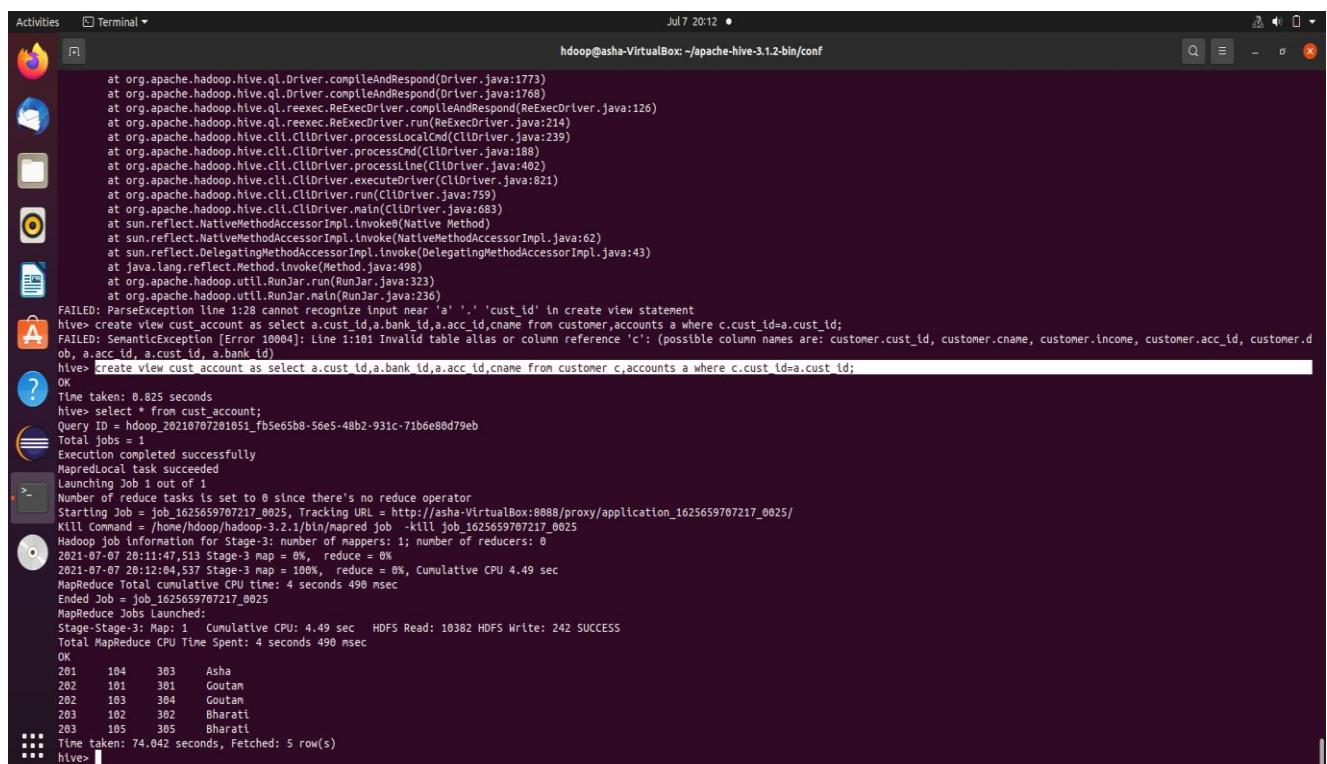


```
Jul 7 20:07 •
hadoop@asha-VirtualBox: ~/apache-hive-3.1.2-bin/conf

at org.apache.hadoop.util.RunJar.run(RunJar.java:323)
at org.apache.hadoop.main(RunJar.java:236)
FAILED: ParseException line 1:56 cannot recognize input near '(' 'between' "'1990-10-10'" in expression specification
hive> select cust_id,cname,dob from customer where dob not between '1990-10-10' and '1993-02-01' group by cust_id,cname,dob order by cust_id DESC;
Query ID = hdoop_20210707200506_86e220b8-841b-4802-82fc-264ecbaef86
Total jobs = 2
Launching Job 1 out of 2
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = Job_1625659707217_0023, Tracking URL = http://asha-VirtualBox:8088/proxy/application_1625659707217_0023/
Kill Command = /home/hadoop/hadoop-3.2.1/bin/mapred job -kill job_1625659707217_0023
2021-07-07 20:05:28,309 Stage-1 map = 0%, reduce = 0%
2021-07-07 20:06:01,672 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 4.35 sec
MapReduce Total cumulative CPU time: 8 seconds 190 msec
Ended Job = job_1625659707217_0023
Launching Job 2 out of 2
Number of reduce tasks determined at compile time: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = Job_1625659707217_0024, Tracking URL = http://asha-VirtualBox:8088/proxy/application_1625659707217_0024/
Kill Command = /home/hadoop/hadoop-3.2.1/bin/mapred job -kill job_1625659707217_0024
2021-07-07 20:07:13,729 Stage-2 map = 100%, reduce = 0%, Cumulative CPU 3.53 sec
2021-07-07 20:07:17,011 Stage-2 map = 100%, reduce = 100%, Cumulative CPU 6.98 sec
MapReduce Total cumulative CPU time: 6 seconds 980 msec
Ended Job = job_1625659707217_0024
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1   Cumulative CPU: 8.19 sec   HDFS Read: 13556 HDFS Write: 210 SUCCESS
Stage-Stage-2: Map: 1 Reduce: 1   Cumulative CPU: 6.98 sec   HDFS Read: 8247 HDFS Write: 221 SUCCESS
Total MapReduce CPU Time Spent: 15 seconds 170 msec
OK
206    Pooja  1995-07-20
204    Pavitra 1990-07-26
202    Goutam  2000-11-16
201    Asha    2000-07-27
Time taken: 133.007 seconds, Fetched: 4 row(s)
hive>
```

Retrieve view which contains all the fields from account and cname in it.

Demonstrate update operation it.



```
Jul 7 20:12 •
hadoop@asha-VirtualBox: ~/apache-hive-3.1.2-bin/conf

at org.apache.hadoop.hive.ql.Driver.compileAndRespond(Driver.java:1773)
at org.apache.hadoop.hive.ql.Driver.compileAndRespond(Driver.java:1768)
at org.apache.hadoop.hive.ql.reexec.ReExeDriver.compileAndRespond(ReExeDriver.java:126)
at org.apache.hadoop.hive.ql.reexec.ReExeDriver.run(ReExeDriver.java:214)
at org.apache.hadoop.hive.cli.Clidriver.processLocalCmd(Clidriver.java:239)
at org.apache.hadoop.hive.cli.Clidriver.processCmd(Clidriver.java:198)
at org.apache.hadoop.hive.cli.Clidriver.processLine(Clidriver.java:402)
at org.apache.hadoop.hive.cli.Clidriver.executeDriver(Clidriver.java:821)
at org.apache.hadoop.hive.cli.Clidriver.run(Clidriver.java:759)
at org.apache.hadoop.hive.cli.Clidriver.main(Clidriver.java:683)
at sun.reflect.NativeMethodAccessorImpl.invoke0(Native Method)
at sun.reflect.NativeMethodAccessorImpl.invoke(NativeMethodAccessorImpl.java:62)
at sun.reflect.DelegatingMethodAccessorImpl.invoke(DelegatingMethodAccessorImpl.java:43)
at java.lang.reflect.Method.invoke(Method.java:498)
at org.apache.hadoop.util.RunJar.run(RunJar.java:223)
at org.apache.hadoop.util.RunJar.main(RunJar.java:236)
FAILED: ParseException line 1:28 cannot recognize input near 'a '.' 'cust_id' in create view statement
hive> create view cust_account as select a.cust_id,a.bank_id,a.acc_id,cname from customer,accounts a where c.cust_id=a.cust_id;
FAILED: SemanticException [Error 10004]: Line 1:101 invalid table alias or column reference 'c': (possible column names are: customer.cust_id, customer.cname, customer.income, customer.acc_id, customer.dob, a.acc_id, a.cust_id, a.bank_id)
hive> create view cust_account as select a.cust_id,a.bank_id,a.acc_id,cname from customer c,accounts a where c.cust_id=a.cust_id;
OK
Time taken: 0.825 seconds
hive> select * from cust_account;
Query ID = hdoop_20210707201051_fb5e65b8-56e5-48b2-931c-71b6e80d79eb
Total jobs = 1
Execution completed successfully
MapredLocal task succeeded
Launching Job 1 out of 1
Number of reduce tasks is set to 0 since there's no reduce operator
Starting Job = Job_1625659707217_0025, Tracking URL = http://asha-VirtualBox:8088/proxy/application_1625659707217_0025/
Kill Command = /home/hadoop/hadoop-3.2.1/bin/mapred job -kill job_1625659707217_0025
Hadoop job information for Stage-3: number of nappers: 1; number of reducers: 0
2021-07-07 20:11:47,513 Stage-3 map = 0%, reduce = 0%
2021-07-07 20:12:04,537 Stage-3 map = 100%, reduce = 0%, Cumulative CPU 4.49 sec
MapReduce Total cumulative CPU time: 4 seconds 490 msec
Ended Job = job_1625659707217_0025
MapReduce Jobs Launched:
Stage-Stage-3: Map: 1   Cumulative CPU: 4.49 sec   HDFS Read: 10382 HDFS Write: 242 SUCCESS
Total MapReduce CPU Time Spent: 4 seconds 490 msec
OK
201    104    303    Asha
202    101    301    Goutam
202    103    304    Goutam
203    102    302    Bharati
203    105    305    Bharati
Time taken: 74.042 seconds, Fetched: 5 row(s)
hive>
```